


Clustering 7.4 – 7.6



- **7.4 Partitioning methods**
- **7.5 Hierarchical methods**
- **7.6 Density-based methods**

1



7.4 Partitioning Methods

- Given a dataset, D , and k , the number of clusters to form, a partitioning algorithm organizes n objects into k partitions, where ($k \leq n$)
- Each partition represents a cluster
- These clusters are formed to optimize an objective partitioning criterion, in terms of attributes
- dissimilarity function (based on distance), where “similar” objects reside within a cluster and “dissimilar” objects reside in other clusters

2



k -Means and k -Medoids

- Heuristic methods: k -means and k -medoids
 - k -means: a cluster is represented by the mean (center) value
 - k -medoids or PAM (partition around medoids): each cluster is represented by one of the objects in the cluster
- Centroid-Based Technique: k -Means Method

3



k -Means and k -Medoids

- Partitions a set of n objects into k clusters,
 - optimizing the *intracluster* similarity
 - minimizing *intercluster* similarity
- Cluster similarity is measured with respect to the mean value of the objects within that cluster → referred to as the cluster's *centroid* or *center of gravity*

4



k -Means algorithm

- Inputs → D – dataset with n objects
- Output → set of k clusters
- Method
 - (1) arbitrarily select k objects from D as the initial cluster mean (centroid)
 - (2) repeat
 - (3) reassign each object to the cluster to which it is most similar, based upon the distance between the object and the cluster mean
 - (4) update cluster mean for objects in each cluster
 - (5) iterate until the criterion function converges (no redistribution of objects in any cluster)
 - known as *iterative relocation*

5



k -Means algorithm

- Use the squared-error criterion – distance from object to centroid is squared and then distances are summed

(below is the *absolute-error criterion* used in *k-Medoids*)

$$E = \sum_{i=1}^k \sum_{p \in C_i} d(p, o_i)^2$$


→ E = sum of square error for all objects

→ p = point in space representing object

→ m_i = mean of cluster, C_i

6

k -Means: Example



Pros and Cons of k -Means

- Relatively efficient – computational complexity: $O(tkn)$
 - n = # objects
 - k = # clusters
 - T = # iterations
 - $k, t \ll n$.



Pros and Cons of k -Means

- Applicable only when mean is defined
 - What about categorical data? – NO (k -modes)
 - Can we normalize occurrence of categorical data? – We might lose some attribute details??
 - Create binary values for categorical data? – We might lose some attribute details??
- Need to specify the number of clusters
- Unable to handle noisy data and outliers

9



Variations of the k -Means

- Aspects
 - Selection of the initial k -means
 - Calculation of dissimilarity
 - Different strategies to calculate cluster means:
 - → one example is to apply a hierarchical agglomerative algorithm, which determines the number of clusters and finds an initial clustering, and then use iterative relocation to improve clustering

10


Variations of the k -Means

- k -modes handle categorical data
 - Use mode instead of mean
 - Mode: the most frequently occurring item
 - A mixture of categorical and numerical data: k -prototype method
- Expectation-Maximum
 - Each object is assigned to each cluster, according to a weight representing the probability of membership in that particular cluster
 - Details → 7.8 Model-Based Clustering Methods

11

Representative Object-Based Technique: k -medoids method

- k -medoids: the most centrally located object in a cluster
- Used because k -means is sensitive to outliers – large values distort the distribution of objects




12

Partitioning Around Medoids (PAM)

- Instead of mean value of objects in the cluster, we arbitrarily choose an object itself, as the initial medoid
- Until no change, do
 - Reassign each object to the cluster to which the nearest medoid *improves* the clustering
 - For each pair of non-select object O_{random} and selected object O_j , calculate the total swapping cost S .
 - If $S < 0$ then swap O_j with O_{random} to form the new set of k -medoids.

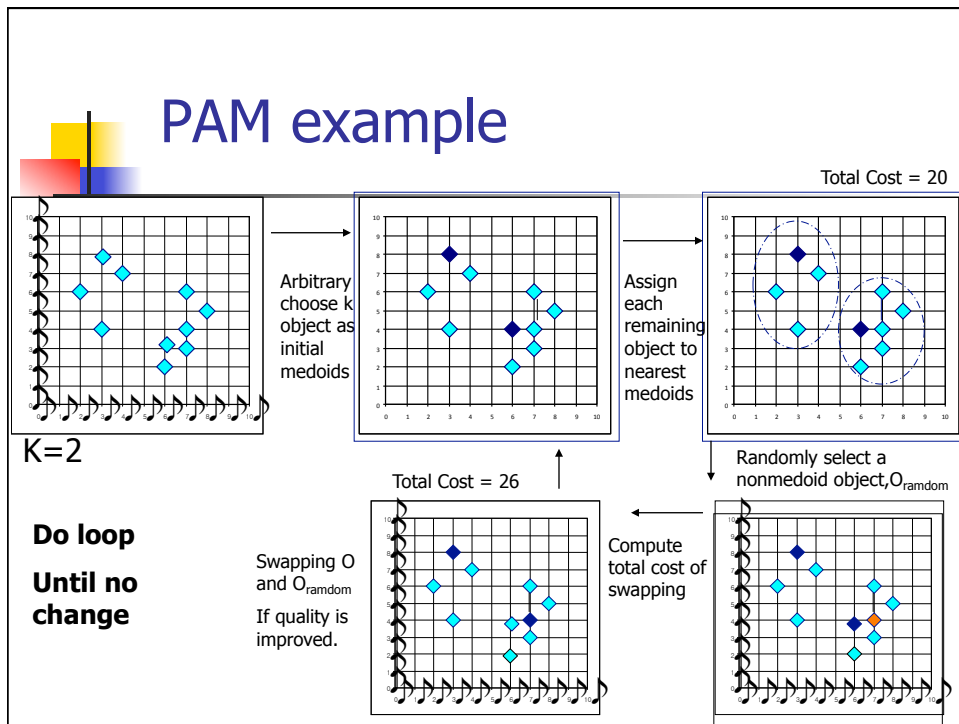
13

k -medoids clustering



- data object
- + cluster center
- before swapping
- after swapping

14



Pros and Cons of PAM

- PAM is more robust than k-means in the presence of noise and outliers
 - Medoids are less influenced by outliers
- However, processing k -Medoids is more costly (in terms of time and memory)
- Both require the user to specify the number of k clusters

16



Pros and Cons of PAM

- We can also use the k-median (middle value) and k-modes (most frequent value)
- PAM is efficient for small data sets but does not scale well for large data sets
 - Iterative computational complexity = $O(k(n-k)^2)$
 - Sampling method = CLARA

17



CLARA – sampling method

- Clustering LARge Applications
- Draw multiple samples (small portion of dataset D) of the data set
- Apply Partitioning Around Medoids on each sample to give the best clustering
- Perform better than PAM in larger data sets
- If selected in a random manner, it should closely represent the original dataset
- Efficiency depends on the sample size
 - A good clustering on samples may not be a good clustering of the whole data set

18



CLARANS

- Clustering Large Applications based upon RANdomized Search
- Since PAM uses entire dataset and CLARA can only select the medoids from the *chosen samples*
- CLARA cannot find the best clustering if any of the best sampled medoids is NOT among the best k-medoids chosen
- Draws a sample with some randomness in each step of the search

19



CLARANS

- Conceptually, clustering process can be seen as a search through a graph with each node as a potential solution
- 2 nodes are neighbors (connected by an arc) if their sets differ by only 1 object
- Each node is can be assigned a cost that is defined by the total dissimilarity of each object and the medoid of its cluster
- Partitioning Around Medoids (PAM) examines all of the current node's neighbors for a minimal cost solution

20

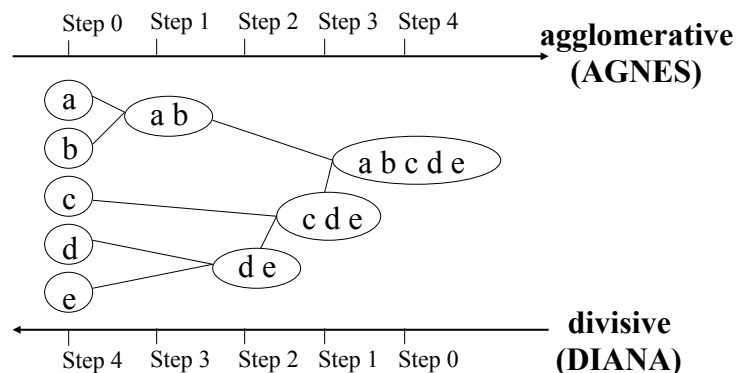
CLARANS (continued)

- CLARA works on a sample – has fewer neighbors to check
- CLARANS *dynamically* draws a random neighbor sample at each step of the process by not confining its search to *local* neighbors
- If a better neighbor, with a lower error is found, the process begins again
- Once the user-specified number of local minima has been found, the solution is output!

21

7.5 Hierarchical Clustering

- Group data objects into a tree of clusters



22


AGNES – Agglomerative NESTing

- Initially, each object is a cluster
- Step-by-step cluster merging, until all objects form a cluster
 - Single-link approach
 - Each cluster is represented by all of the objects in the cluster
 - The similarity between two clusters is measured by the similarity of the closest pair of data points belonging to different clusters


23

DIANA – Divisive ANALysis

- Initially, all objects are in one cluster
- Step-by-step splitting clusters until each cluster contains only one object
 - Split according to some principle: the maximum Euclidean distance between the closest neighboring objects in the cluster.



24



- Dendrogram → Tree structure to represent the process of hierarchical clustering
- Shows how to merge clusters hierarchically
- Decompose data objects into a multi-level nested partitioning (a tree of clusters)
- A clustering of the data objects: cutting the dendrogram at the desired level
 - Each connected component forms a cluster
 - $\{\{\{A,B\},C\},\{\{\{D,E\},F\},G\},\{H,I\}\}$

25

Distance Measures in Dendrogram

- Single link: Minimum distance – nearest-neighbor clustering

$$d_{\min}(C_i, C_j) = \min_{p \in C_i, q \in C_j} d(p, q)$$
 - Complete link: Maximum distance – farthest-neighbor clustering

$$d_{\max}(C_i, C_j) = \max_{p \in C_i, q \in C_j} d(p, q)$$
 - Centroid: Mean distance

$$d_{\text{mean}}(C_i, C_j) = d(m_i, m_j)$$
 - Average link: Average distance

$$d_{\text{avg}}(C_i, C_j) = \frac{1}{n_i n_j} \sum_{p \in C_i} \sum_{q \in C_j} d(p, q)$$
- m: mkan for a cluster
 → C: a cluster
 → n: the number of objects in a cluster

26



Hierarchical Clustering Challenges

- Hard to choose merge/split points
 - Cannot undo merging/splitting
 - Merging/splitting decisions are critical
- Computational complexity is high: $O(n^2)$
- Integrating hierarchical clustering with other techniques
 - BIRCH, ROCK, CHAMELEON

27



BIRCH - for numeric attributes

- **B**alanced **I**terative **R**educing and **C**lustering using **H**ierarchies
- Overcomes two difficulties of agglomerative clustering:
 - scalability
 - The inability to undo what was done in the previous step
- Introduces *clustering feature* and *clustering feature tree*
- CF (Clustering Feature) tree: a 3-D hierarchical data structure summarizing object information
 - Clustering objects → clustering leaf nodes of the CF tree
 - Summary of statistics for a given cluster

28



Clustering Feature Vector


Clustering Feature: $CF = (N, LS, SS)$

(data points) N : **Number of data points**

(Linear Sum) LS : $\sum_{i=1}^N X_i$

(Square Sum) SS : $\sum_{i=1}^N X_i^2$

$CF = (5, (16,30),(54,190))$



(3, 4)
(2, 6)
(4, 5)
(4, 7)
(3, 8)

29



Clustering Feature Vector

- These features are *additive*
- If CF1 has points (2,5), (3,2), (4,3)
- $CF1 = \{3, (2+3+4, 5+2+3), (2^2 + 3^2 + 4^2, 5^2 + 2^2 + 3^2)\} = \{3, (9,10), (29,38)\}$
- Let $CF2 = \{3, (35,36), (417,440)\}$
- We add these to get $\rightarrow CF3 = \{6, (44,46), (446,478)\}$

30



BIRCH

- Phase 1:
 - scans database to build an initial in-memory tree, which tries to preserve the inherent clustering structure of the data
 - An object is inserted into the closest leaf entry.
 - If the diameter of the subcluster stored in the leaf node after insertion is larger than the threshold value, then the leaf node are split.
- Phase 2:
 - applies a selected clustering algorithm to cluster the leaf nodes by removing sparse clusters as outliers and grouping the dense clusters

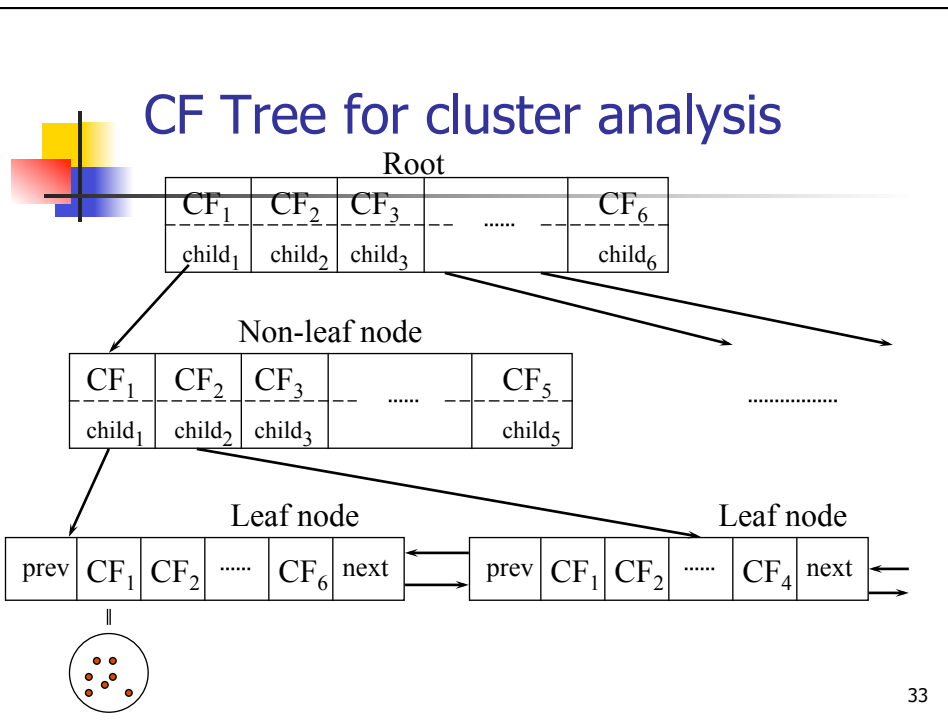
31



CF Tree in BIRCH

- Clustering feature:
 - Summarize the statistics for a sub-cluster
 - Register computing cluster measurements and efficiently utilizes storage space
- A CF tree: a height-balanced tree storing the clustering features for a hierarchical clustering
 - A non-leaf node in a tree has children
 - The non-leaf nodes store sums of the CF's of children

32



- ## CF Tree info
- Built dynamically as new data are inserted
 - Used to guide a new insertion into the correct cluster
 - A height-balanced tree
 - 2 parameters:
 - branching factor B = max num of children for non-leaf node and
 - threshold T for = max diameter (or radius) of sub-clusters at the leaf nodes
 - Each non-leaf node contains at most B entries
 - Each leaf node has "prev" and "next" pointers to chain all leaf nodes together
 - Diameter (or radius) of all entries in a leaf node must be less than T
- 34



Pros and Cons of BIRCH

- Linear scalability
 - Good clustering with a single scan
- Can handle only numeric data
- Sensitive to the order of the data records
- Computational complexity of $O(n)$

35



Clustering Categorical Data: The ROCK Algorithm

- ROCK: **RO**bust **C**lustering using **linK**s
 - S. Guha, R. Rastogi & K. Shim, ICDE' 99
- Basic Ideas:
 - If two data objects have similar neighborhoods, then the two data objects belong to the same cluster, and thus can be merged.
 - Use links (common neighbors between two objects) to measure similarity/proximity
 - Not distance-based

36

Similarity Measure in ROCK

- Traditional measures for categorical data may not work well, e.g., Jaccard coefficient
- Example: Two groups (clusters) of transactions
 - C_1 . $\langle a, b, c, d, e \rangle$: $\{a, b, c\}, \{a, b, d\}, \{a, b, e\}, \{a, c, d\}, \{a, c, e\}, \{a, d, e\}, \{b, c, d\}, \{b, c, e\}, \{b, d, e\}, \{c, d, e\}$
 - C_2 . $\langle a, b, f, g \rangle$: $\{a, b, f\}, \{a, b, g\}, \{a, f, g\}, \{b, f, g\}$
- Jaccard co-efficient-based similarity function:

$$Sim(T_1, T_2) = \frac{|T_1 \cap T_2|}{|T_1 \cup T_2|}$$
 - Ex. Let $T_1 = \{a, b, c\}, T_2 = \{c, d, e\}$

$$Sim(T_1, T_2) = \frac{|\{c\}|}{|\{a, b, c, d, e\}|} = \frac{1}{5} = 0.2$$
- Jaccard co-efficient may lead to wrong clustering result
 - C_1 : 0.2 ($\{a, b, c\}, \{b, d, e\}$) to 0.5 ($\{a, b, c\}, \{a, b, d\}$)
 - C_1 & C_2 : could be as high as 0.5 ($\{a, b, c\}, \{a, b, f\}$)

37

Link Measure in ROCK

- Links: # of common neighbors
- C_1 $\langle a, b, c, d, e \rangle$: $\{a, b, c\}, \{a, b, d\}, \{a, b, e\}, \{a, c, d\}, \{a, c, e\}, \{a, d, e\}, \{b, c, d\}, \{b, c, e\}, \{b, d, e\}, \{c, d, e\}$
 - C_2 $\langle a, b, f, g \rangle$: $\{a, b, f\}, \{a, b, g\}, \{a, f, g\}, \{b, f, g\}$
- Assume $\theta = 0.5$ (use it to find neighbors)
- Let $T_1 = \{a, b, c\}, T_2 = \{c, d, e\}, T_3 = \{a, b, f\}$
 - $link(T_1, T_2) = 4$, since they have 4 common neighbors
 - $\{a, c, d\}, \{a, c, e\}, \{b, c, d\}, \{b, c, e\}$
 - $link(T_1, T_3) = 3$, since they have 3 common neighbors
 - $\{a, b, d\}, \{a, b, e\}, \{a, b, g\}$
- Thus link is a better measure than Jaccard coefficient

38

7.6 Density-Based Clustering Methods

- Clustering based on density (local cluster criterion), such as density-connected points
- Major features:
 - Discover clusters of arbitrary shape
 - Handle noise
 - One scan
 - Need density parameters as termination condition


39

Density-Based Clustering: Basic Concepts

- Two parameters:

- **Eps**: Maximum radius of the neighbourhood
- **MinPts**: Minimum number of points in an Eps-neighbourhood of that point
- $N_{Eps}(p)$: $\{q \text{ belongs to } D \mid \text{dist}(p,q) \leq Eps\}$
- **Directly density-reachable**: A point p is directly density-reachable from a point q w.r.t. Eps , $MinPts$ if
 - p belongs to $N_{Eps}(q)$
 - core point condition:

$$|N_{Eps}(q)| \geq MinPts$$





MinPts = 5
Eps = 1 cm

40

Density-Reachable and Density-Connected


- Density-reachable:
 - A point p is **density-reachable** from a point q w.r.t. Eps , $MinPts$ if there is a chain of points p_1, \dots, p_n , $p_1 = q$, $p_n = p$ such that p_{i+1} is directly density-reachable from p_i
- Density-connected
 - A point p is **density-connected** to a point q w.r.t. Eps , $MinPts$ if there is a point o such that both, p and q are density-reachable from o w.r.t. Eps and $MinPts$



41

DBSCAN: Density Based Spatial Clustering of Applications with Noise

- Relies on a *density-based* notion of cluster: A *cluster* is defined as a maximal set of density-connected points
- Discovers clusters of arbitrary shape in spatial databases with noise



42



DBSCAN: The Algorithm

- Arbitrary select a point p
 - Retrieve all points density-reachable from p w.r.t. Eps and $MinPts$.
 - If p is a core point, a cluster is formed.
 - If p is a border point, no points are density-reachable from p and DBSCAN visits the next point of the database.
 - Continue the process until all of the points have been processed.