


First-Come, First-Served (FCFS) Scheduling

Process Burst Time

P_1 24
 P_2 3
 P_3 3

- Suppose that the processes arrive in the order: P_1, P_2, P_3

The Gantt Chart for the schedule is:




- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

P_2, P_3, P_1

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- Convoy effect short process behind long process

Shortest-Job-First (SJF) Scheduling


- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
 - nonpreemptive – once CPU given to the process it cannot be preempted until completes its CPU burst
 - preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is optimal – gives minimum average waiting time for a given set of processes

Example of Non-Preemptive SJF

Process Arrival Time Burst Time

P_1 0.0 7
 P_2 2.0 4
 P_3 4.0 1
 P_4 5.0 4

- SJF (non-preemptive)




- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

Process Arrival Time Burst Time

P_1 0.0 7
 P_2 2.0 4
 P_3 4.0 1
 P_4 5.0 4

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process


Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high

Example of RR with Time Quantum = 20


Process	Burst Time
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better *response*

Time Quantum and Context Switch Time




Turnaround Time Varies With The Time Quantum



Multilevel Queue

- Ready queue is partitioned into separate queues: foreground (interactive) background (batch)
- Each queue has its own scheduling algorithm
 - foreground – RR
 - background – FCFS
- Scheduling must be done between the queues
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS

Multilevel Queue Scheduling




Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

Example of Multilevel Feedback Queue

- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 . When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Multilevel Feedback Queues



Multiple-Processor Scheduling

- CPU scheduling more complex when multiple CPUs are available
- *Homogeneous processors* within a multiprocessor
- *Load sharing*
- *Asymmetric multiprocessing* – only one processor accesses the system data structures, alleviating the need for data sharing

Real-Time Scheduling

- *Hard real-time* systems – required to complete a critical task within a guaranteed amount of time
- *Soft real-time* computing – requires that critical processes receive priority over less fortunate ones

Pthread Scheduling API

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_t init(&attr);
    /* set the scheduling algorithm to PROCESS or SYSTEM */
    pthread_attr_t setscope(&attr, PTHREAD_SCOPE_SYSTEM);
    /* set the scheduling policy - FIFO, RT, or OTHER */
    pthread_attr_t setschedpolicy(&attr, SCHED_OTHER);
    /* create the threads */
    for (i = 0; i < NUM_THREADS; i++)
        pthread_create(&tid[i], &attr, runner, NULL);
}
```

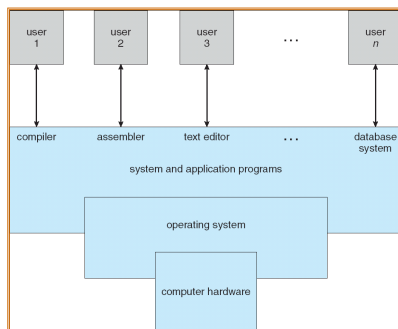
Pthread Scheduling API

```
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}
/* Each thread will begin control in this
function */
void *runner(void *param)
{
    printf("I am a thread\n");
    pthread_exit(0);
}
```


Review Chapter 1

- Overview of the “Operating System”
- Kernel mode/user mode*
- Storage (I/O)
- Memory
- Interrupt *

Four Components of a Computer System



Storage-Device Hierarchy



Chapter 2

- How to design an operating system
- Structure of an “operating system” *
 - Layers
 - Microkernels
 - Modules
 - Comparing them?
- What’s a system call? *
- Eg. Question: What are the advantages and disadvantages of using the same system call interface for manipulating both files and devices?
- Eg. Question: in what ways is the modular kernel approach similar to the layered one? In what ways they’re different?
- Virtual machine?

Chapter 3 Review

- Process Concept *
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Interprocess Communication*
- Communication in Client-Server Systems
- RPC (*stub, marshal*)
- RMI

Chapter 3 Review

- Process Scheduling
 - Status diagram
 - I/O-bound process and a CPU-bound process
 - Queues (*many*)
 - Context switch
 - Process life cycle (create, terminate)

Chapter 3 Review

- Interprocess Communication
 - Shared memory
 - Producer-consumer problem
 - Data structure (*circular array*)
 - Message passing
 - Direct/indirect
 - synchronization

Chapter 3 Review

- Communication in Client-Server Systems/RPC/RMI
 - What a socket? Port number?
 - RPC: port? Marshalling?
 - RMI: concept of stub, skeleton, marshalling
 - Compare RPC/RMI

Chapter 4 Threads

- Multithreading Models*
 - 1 to 1
 - 1 to many
 - Many to many
- Threading Issues
 - Fork() and exec()
- PThread/Win32/Java Threading
- Compare thread and process
 - Eg. Q: which of the following are shared across threads...
 - Register
 - Heap
 - Globals
 - Stack

Chapter 5

- Basic Concepts
- Scheduling Criteria
 - Turnaround time
 - Waiting time
 - Response time ...
- Scheduling Algorithms *
- FCFS
- SJF
- RR
- Priority ...
- Multiple-Processor Scheduling
- Eg. *The exercise handout*